Brian Farrell

Big Data Exercise 2

1. [15 points] Find the column names in the Opioid dataset. The naive way is to gunzip the .gz file and run head -1 on the result, but you likely don't have enough disk space. Conveniently, zcat can read the file and write the unzipped contents into stdout, which can be piped into head -1.
    a. I used the following command to answer this question: gzcat arcos_all_washpost.tsv.gz | head -1
        i. Gzcat decompresses and outputs the content of the arcos_all_washpost.tsv.gsv file. I used gzcat because I ran this command on my mac laptop and gzact is the mac version of zcat. Head -1 outputs the header row of the data and resulted in the below output.
    b. REPORTER_DEA_NO REPORTER_BUS_ACT REPORTER_NAME REPORTER_ADDL_CO_INFO REPORTER_ADDRESS1 REPORTER_ADDRESS2 REPORTER_CITY REPORTER_STATE REPORTER_ZIP REPORTER_COUNTY BUYER_DEA_NO BUYER_BUS_ACT BUYER_NAME BUYER_ADDL_CO_INFO BUYER_ADDRESS1 BUYER_ADDRESS2 BUYER_CITY BUYER_STATE BUYER_ZIP BUYER_COUNTY TRANSACTION_CODE DRUG_CODE NDC_NO DRUG_NAME QUANTITY UNIT ACTION_INDICATOR ORDER_FORM_NO CORRECTION_NO STRENGTH TRANSACTION_DATE CALC_BASE_WT_IN_GM DOSAGE_UNIT TRANSACTION_ID

Product_Name Ingredient_Name Measure MME_Conversion_Factor

Combined_Labeler_Name Revised_Company_Name Reporter_family dos_str

2. [15 points] Find the number of rows in the Opioid dataset by processing the zcat output, stripping the header row, and counting the remaining lines using wc.
    a. I ran the following command to find the number of rows minus the headers:
    b. zcat arcos_all_washpost.tsv.gz | awk 'NR > 1' | wc -l
    c. The resulting output of the above command was: 178,598,026.
    d. I used zcat to look into the unzipped file. Awk ‘NR > 1’ was used to skip the header row. And finally, wc -1 was used to count up the remaining number of rows in the file.
    e. I switched to my windows pc on this question so my commands would run faster, which is why I switched from gzcat to zcat.
3. [20 points] Find the names of all the drugs named in the dataset.
    a. I used the following python program to print all the unique drug names into a text file called “drug_names.txt”

```
import gzip


def get_drugs():
   drug_names = set()  # Set to avoid duplicates
   with gzip.open('arcos_all_washpost.tsv.gz', 'rt') as f:
      header = f.readline()
```

```
        drug_column_index = header.strip().split('\t').index('DRUG_NAME')  # Find the
index for the drug name column
        for line in f:
            drug_names.add(line.strip().split('\t')[drug_column_index])  # Add drug
names to the set
    return drug_names


def write_drugs(drug_names):
    with open('drug_names.txt', 'w') as f:
        for drug in (drug_names):  # for loop of the set the set
            f.write(drug + '\n')  # Write each drug name to the file


drug_names = get_drugs()
write_drugs(drug_names)
```

b. The output of the text file was:
    i. HYDROCODONE
    ii. OXYCODONE

4. [20 points] Estimate the number of rows for each year in the dataset. There may be enough space in the shell, but this exercise requires you to assume that that's not the case. So here's a potential strategy: Use the shuf command to extract, say, random 7,500 rows

from the output of zcat. Find the proportion of rows for each year in this extract. Assuming that the distribution of the random 7,500 rows is similar to the distribution in the whole file, estimate the number of rows for each year.

a. I used the following python program to estimate the number of rows in each year.

```
import gzip
import random
from datetime import datetime


def estimate_years(file_path, sample_size=7500, total_rows=178598026):
    with gzip.open(file_path, 'rt') as f:
        header = f.readline()
        columns = header.strip().split('\t')
        transaction_date_index = columns.index('TRANSACTION_DATE')  # Find the column index
        # Extract a random sample of rows
        sample = []
        for i, line in enumerate(f):
            if random.random() < sample_size / total_rows:
                fields = line.strip().split('\t')
```

```
            if len(fields) > transaction_date_index:  # Ensure the column exists
                date_str = fields[transaction_date_index]  # Extract the date
                # print(f"Processing row {i + 1}: Date = {date_str}")  # Debug: Print the date being
processed
                try:
                    # Parse the date in "Month/Day/Year" format and extract the year
                    date = datetime.strptime(date_str, '%m%d%Y')  #formatting the date
                    year = date.year
                    sample.append(year)
                except ValueError:
                    # Skipping null dates
                    print(f"Skipping invalid date: {date_str}")  # Debug: Print invalid dates
                    continue
            else:
                print(f"Skipping row {i + 1}: Missing TRANSACTION_DATE") #Used to debug
rows
            if len(sample) >= sample_size:
                break
```

```
    # Count the occurrences of each year in the sample
    year_counts = {}
    for year in sample:
        year_counts[year] = year_counts.get(year, 0) + 1

    # Step 3: Estimate the rows
    scaling_factor = total_rows / sample_size
    estimated_counts = {year: count * scaling_factor for year, count in year_counts.items()}

    # Print answer
    print("\nYear\tEstimated Rows")
    for year, count in sorted(estimated_counts.items()):
        print(f"{year}\t{count:.0f}")

if __name__ == "__main__":
    file_path = 'arcos_all_washpost.tsv.gz'  # Path to the dataset (I had trouble at first so I had to navigate straight to the directory)
    estimate_years(file_path)
```

b. The output of this file was:

i. Year Estimated Rows
ii. 2006 19741035
iii. 2007 23670192
iv. 2008 24432210
v. 2009 26003873
vi. 2010 27194526
vii. 2011 28623310
viii. 2012 28932880

c. I averaged the estimated rows and got 25,514,004 rows per year according to this sample of 7500 entries.

d. I attempted this first in shell, but kept running into issues. I have more experience writing python scripts, so I switched to it because I felt more comfortable debugging python.

5. [15 points] Obtain the count of rows for June 2012 by extracting all such rows from arcos_all_washpost.tsv.gz and running wc on the extracted rows.

   a. I ran the command: zgrep -P '\t06[0-9]{2}2012\t' arcos_all_washpost.tsv.gz | wc -l

      i. The result was: 2,323,389

6. [15 points] Estimate the count of rows for June 2012 based on answers to questions 1, 2, and 4 and compare that count with your findings from Q5.

   a. There are around 25,514,004 rows per year according to my answer to question four. Dividing this by 12, to get the monthly average, equals 2,126,167 rows for June 2012. This is around 200,000 off the actual number of rows in June 2012.

This is a pretty good estimate considering the data used was only a 7500 random sample of the 178 million rows.